

[Team 42] ProjF4 Final Report: LSTM and CNN Based Stock Market Predictor

Alexander Kyu
Email: awkyu@ncsu.edu

Anthony Kyu
Email: amkyu@ncsu.edu

Michael McKnight
Email: mpmcknig@ncsu.edu

I. ABSTRACT

Predicting the stock market is a challenging task that has been attempted with a variety of methods. In this project, a neural network model that combines aspects of both CNNs and LSTMs to predict the appropriate position of a company's stock (Buy, Sell, or Hold) using stock data and technical indicators commonly used by analysts. The proposed model outperformed the baseline model, but still yielded unfavorable results.

II. MOTIVATION – MAKE DUMMY THICCCCK STACKS

Predicting the stock market has always been a classical neural network problem. However, very few, if any, models (CNN or LSTM) have been very successful in predicting outcomes of the stock market accurately [1]. Therefore, we propose a new neural network model based off literature that combines a convolutional neural network with a LSTM neural network model in hopes to surpass the results of current models [2].

We want to build a model which will be able to predict appropriate stock market positions (Buy, Sell, or Hold) for specific ticker symbols (Advanced Micro Devices -AMD) using various key statistics and market indicators. These are essentially extracted features from basic stock data such as opening and closing prices. With stocks being driven primarily by investor movement and business cycles, it may be increasingly important to understand if these movements can be captured by these market indicators and statistics.

As a problem statement, many current models only use basic stock data, which may lack crucial features for prediction. A model that uses statistics that captures changes over time may improve the accuracy of a Neural Network. A technical difficulty that could arise is the inherent randomness of the stock market. Because news is unpredictable and somewhat random and people are driven to buy shares based on news, these indicators could possibly be less accurate in predicting trends.

III. DATA

The data was collected from Yahoo Finance's historical stock data, which includes open price, high, low, closing price,

adjusted closing price, and volume [3]. The data was further augmented by creating over 150 technical indicators using the TA-Lib Library.

Pattern Recognition		Momentum Indicators	
COL2CROWS	Two Crows	ADX	Average Directional Movement Index
COL3BLACKCROWS	Three Black Crows	ADXR	Average Directional Movement Index Rating
COL3INSIDE	Three Inside Up/Down	ARO	Aroon
COL3LINESTRIKE	Three-Line Strike	AROON	Aroon Oscillator
COL3OUTSIDE	Three Outside Up/Down	AROOMOSC	Balance of Power
COL1STARSINSOUTH	Three Stars In The South	BOP	Commodity Channel Index
COL2MIDSTARSINSOUTH	Three Advancing White Soldiers	CCI	Chande Momentum Oscillator
COLABANDONDORNBAY	Abandoned Baby	CMO	Directional Movement Index
COLADVANCEBLOCK	Advance Block	DM	Moving Average Convergence/Divergence
COLBELTHOLD	Belt-hold	MACD	MACD with controllable MA type
COLBREAKAWAY	Breakaway	MACDXT	Moving Average Convergence/Divergence Fix 12/26
COLCLOSINGMARUBOZU	Closing Marubozu	MACDXTX	Money Flow Index
COLCONCEALBABYSWALL	Concealing Baby Swallow	MF1	Minus Directional Indicator
COLCOUNTERATTACK	Counterattack	MINUS_DI	Minus Directional Movement
COLDARKCLOUDCOVER	Dark cloud cover	MINUS_DM	Momentum
COLDOJI	Doji	PLUS_DI	Plus Directional Indicator
COLDOJISTAR	Doji Star	PLUS_DM	Plus Directional Movement
COLDRAGONFLYDOJI	Dragonfly Doji	PPO	Percentage Price Oscillator
COLENGULFING	Engulfing Pattern	RCC	Rate of change 1: (price/prevprice)-1)*100
COLERVENINGDOJISTAR	Evening Doji Star	ROC	Rate of change Percentage: (price-prevprice)/prevprice
COLERVENINGSTAR	Evening Star	ROCP	Rate of change ratio 100 scale: (price/prevprice)*100
COLGAPSIDESIDEWHITE	Up/Down-gap side-by-side white lines	ROCK100	Relative Strength Index
COLGRAVESTONEDOJI	Gravestone Doji	RSI	Stochastic
COLHAMMER	Hammer	STOCH	Stochastic Fast
COLHANGINGMAN	Hanging Man	STOCHF	Stochastic Relative Strength Index
COLHARAMI	Harami Pattern	STOCHRSI	1-day Rate-of-Change (ROC) of a Triple Smooth EMA
COLHARAMICROSS	Harami Cross Pattern	ULTOSC	Ultimate Oscillator
COLHICKKAKE	Hikkake Pattern	WILLR	Williams' %R
COLHICKKAKEMOD	Modified Hikkake Pattern	Overlap Studies	
COLHOMINGPIGEON	Homing Pigeon	BBANDS	Bollinger Bands
COLIDENTICAL3CROWS	Identical Three Crows	DEMA	Double Exponential Moving Average
COLINNECK	In-neck Pattern	EMA	Exponential Moving Average
COLINVERTEDHAMMER	Inverted hammer	HT_TRENDLINE	Hilbert Transform - Instantaneous Trendline
COLKICKING	Kicking	KAMA	Kaufman Adaptive Moving Average
COLKICKINGBULLLENGTH	Kicking - bull/bear determined by the longer marubozu	MA	Moving average
COLLADDERBUTTON	Ladder Bottom	MAHA	MESA Adaptive Moving Average
COLLONGLEGGEDDOJI	Long Legged Doji	MAVP	Moving average with variable period
COLLONGLINE	Long Line Candle	MDPOINT	Midpoint over period
COLMARUBOZU	Marubozu	MDPRICE	Midpoint Price over period
COLMATHOLD	Math Hold	SAR	Parabolic SAR
COLMORNINGDOJISTAR	Morning Doji Star	SAREXT	Parabolic SAR - extended
COLMORNINGSTAR	Morning Star	SMA	Simple Moving Average
COLONNECK	On-neck Pattern	TS	Triple Exponential Moving Average (TS)
COLPIERCING	Piercing Pattern	TEMA	Triple Exponential Moving Average
COLRICKSHAWMAN	Rickshaw Man	TRIMA	Triangular Moving Average
COLRISEFALLTHREEMETHODS	Rising/falling Three Methods	WMA	Weighted Moving Average
COLSEPARATINGLINES	Separating Lines	Volume Indicators	
COLSHORTLINE	Short Line Candle	AO	Chalkin A/D Line
COLSPINNINGTOP	Spinning Top	AOSC	Chalkin A/D Oscillator
COLSTALLPATTERN	Stalled Pattern	OBV	On Balance Volume
COLSTICKSANDWICH	Stick Sandwich	Cycle Indicators	
COLTASUKI	Tasuki (Dragonfly doji with very long lower shadow)	HT_DPERIOD	Hilbert Transform - Dominant Cycle Period
COLTHRUSTING	Thrusting Pattern	HT_DPHASE	Hilbert Transform - Dominant Cycle Phase
COLTRISTAR	Tristar Pattern	HT_PHASOR	Hilbert Transform - Phasor Components
COLUNIQUE3RIVER	Unique 3 River	HT_SINE	Hilbert Transform - Sinewave
COLUPSIDEGAP2CROWS	Upside Gap Two Crows	HT_TRENDMODE	Hilbert Transform - Trend vs Cycle Mode
COLUPSIDEGAP3METHODS	Upside/Downside Gap Three Methods	Price Transform	
Statistic Functions		AVPRICE	Average Price
BETA	Beta	MDPRICE	Median Price
CORREL	Pearson's Correlation Coefficient (r)	TRPRICE	Typical Price
LINEARREG	Linear Regression	WCLPRICE	Weighted Close Price
LINEARREG_ANGLE	Linear Regression Angle	Volatility Indicators	
LINEARREG_INTERCEPT	Linear Regression Intercept	ATR	Average True Range
LINEARREG_SLOPE	Linear Regression Slope	NATR	Normalized Average True Range
STDEV	Standard Deviation	TRANGE	True Range
TSF	Time Series Forecast	Data From Yahoo Finance	
VAR	Variance	Open	Open
		Close	Close
		High	High
		Low	Low
		Adjusted closed	Adjusted closed
		Volume	Volume

Figure 1. Data, Features, and Technical Indicators Used in Models

The indicators fall within the following groups: overlap studies, momentum indicators, volume indicators, volatility indicators, price transform, cycle indicators, and pattern recognition. The indicators were separated depending on if they were periodic or non-periodic. The periodic indicators were calculated over a range of periods from 5 to 20. The periodic data was then arranged into images [4], where the row is the

period, and the columns are indicators. These images were either then used in the baseline CNN or in the proposed CNN AE.

Due to the overall upward trend for AMD the hold label dominated the classes. To combat this, we down sampled this dominant output to have a more general model that could translate to stocks that are less consistently growing or are in a downward trend.

The predicted data was classified into Sell (0), Buy (1), or Hold (2) based on the closing price within a window of 11 days. Window maximums were labelled as the sell class, window minimums were labeled as the buy class, and all other days were labeled as the hold class [4].

IV. METHODOLOGY

This section outlines the different and altered models that we used to solve this problem of predicting stock market prices. We used a basic 2D CNN as our baseline model and implemented and incorporated our other proposed model to see how those variations change the output and loss of our model. In both models, several dropout layers were used to prevent overfitting and the final layer was a softmax dense layer with three neurons, one for each class.

A. Baseline

As stated above, a 2D Convolutional Neural Network was created to predict stock market prices. This consisted of several layers of convolution and max pooling, followed by a flattening and dense layer. The images of the periodic data were used as the inputs to this model and all non-periodic data was ignored. For each image, the rows were the period window size, and the columns were the technical indicator (Figure 1).

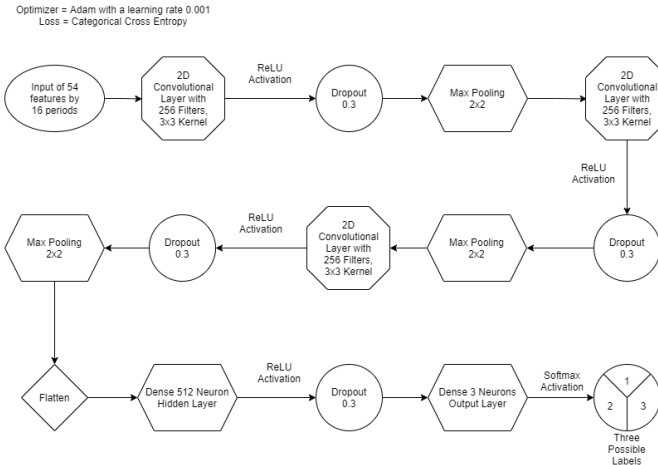


Figure 1: Baseline CNN Model Architecture

For optimizing the baseline model, the Talos library was used to determine the best batch size, convolutional layer neurons, dense layer neurons, dropout, learning rate, and activation function. To reduce computation time a random search of 10%

of all the possible combinations was performed and those randomly chosen combinations were used to train separate models for comparison. In total 32 separate models were trained to determine the optimal parameters (Figure 2).

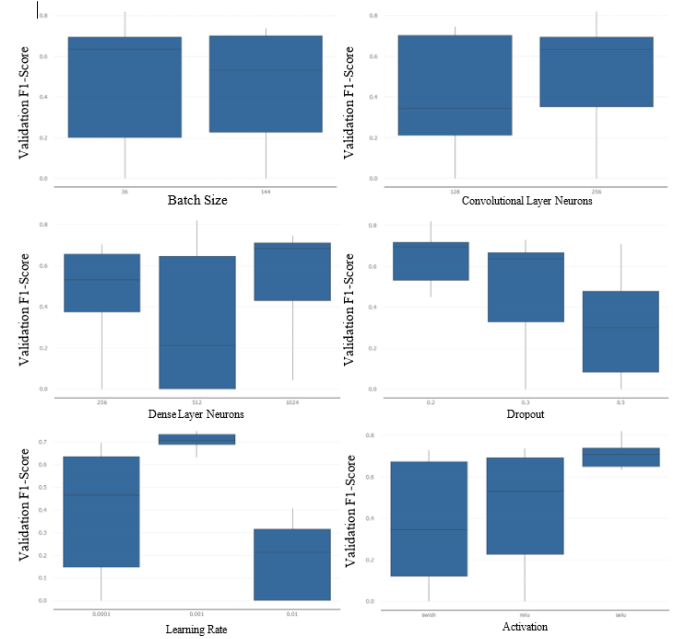


Figure 2: Summary of Talos Optimization for Baseline Model

B. Proposed

We extended our baseline CNN model by converting the CNN to a CNN Autoencoder (AE) to extract features from the periodic data. These encoded features produced from the encoder component of the AE were then flattened and then appended to the nonperiodic data and were used as inputs to a Long Short Term Memory Neural Network (LSTM NN) (Figure 3).

For optimizing the proposed model, the Talos library was used on the autoencoder and LSTM portion separately. For the autoencoder, the parameters optimized were number of kernels, learning rate, kernel size, activation function, and batch size. To reduce computation time a random search of 20% of all the possible combinations was performed and those randomly chosen combinations were used to train separate models for comparison. In total 43 separate models were trained to determine the optimal parameters for the autoencoder. For the LSTM, the parameters optimized were number of nodes, learning rate, dropout, recurrent dropout, and batch size. A random search of 10% of all the possible combinations was performed and resulted in 10 separate models for comparison. After the optimization search (Figure 4) we found for the autoencoder using ReLU, 128 kernels, a kernel size of 5x5, a learning rate of 0.001, and a batch size of 32 was best. We found the best parameters for the LSTM were 256 nodes, a learning rate of 0.001, a dropout of 0.3, a recurrent dropout of 0.2, and a batchsize of 32 (Figure 5).

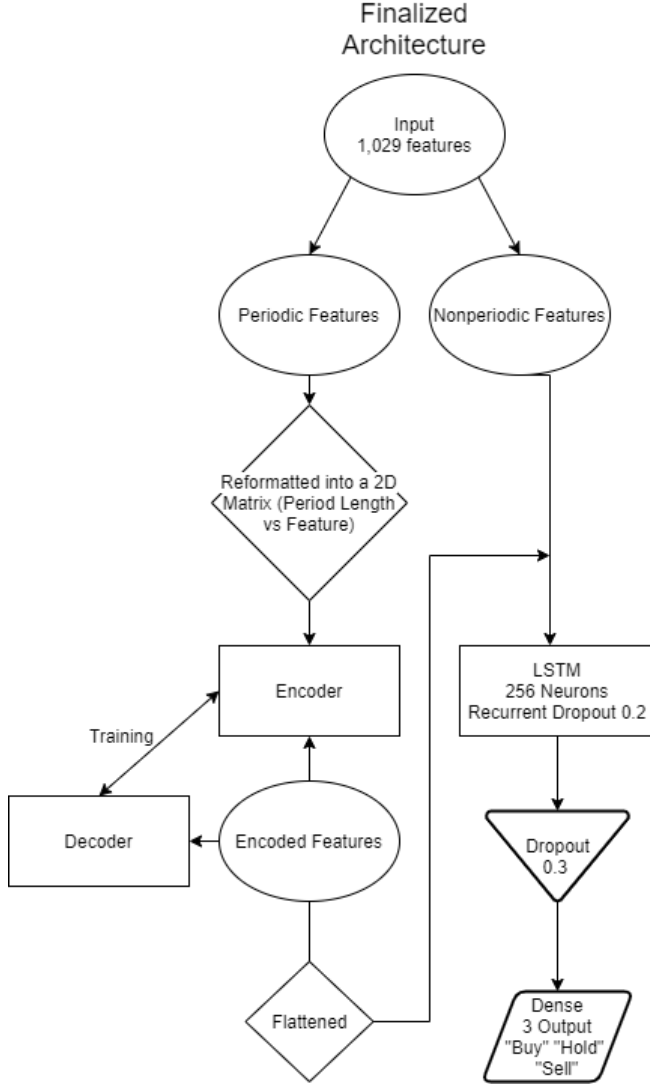


Figure 3: Proposed Model Architecture

V. RESULTS & EVALUATION – MAKE THINNNNN STACKS

A. Computational Model Evaluation

The accuracy per class of the proposed and baseline models was determined by the diagonal of a confusion matrix. The baseline model did not show a preference for any one class but had an overall worse accuracy than the proposed model (Table 1). The proposed model did not label any data points class 0 (sell) or class 1 (buy) despite us under sampling to handle the data imbalance (Table 2).

B. Financial Model Evaluation

A financial evaluation was performed on the last year (April 24th, 2020 to April 23rd, 2021) of AMD's stock closing prices. Our model predicted the appropriate positions of the next day's stock price. These predictions were then used to determine when

stocks were bought, sold and held. If no stock was bought yet, then a Buy position would simulate buying the stock, and a Sell

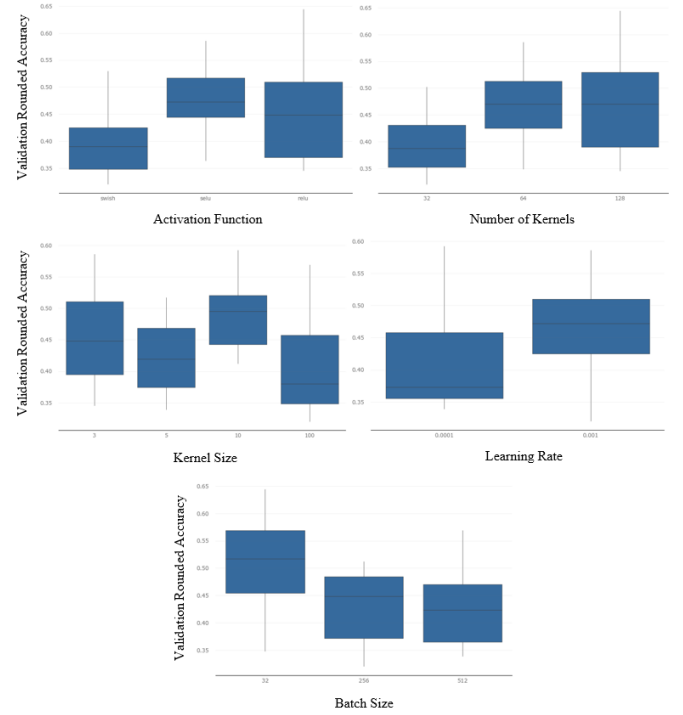


Figure 4: Summary of Talos Optimization for Autoencoder of Proposed Model

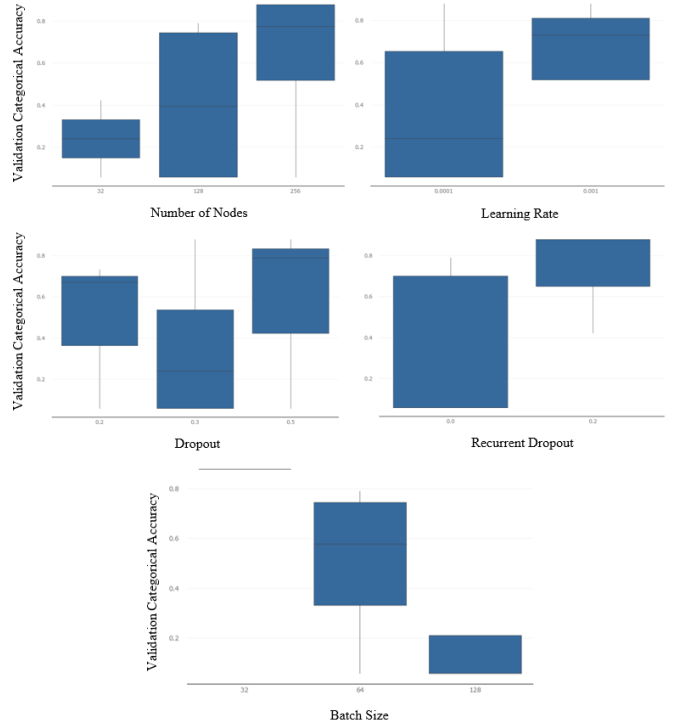


Figure 5: Summary of Talos Optimization for LSTM of Proposed Model

or Hold position would do nothing. If stock was already bought, then a Sell position would simulate selling the stock, while a Buy or Hold position would do nothing. At the end of the year, any stocks held would automatically be sold, and the annual gain was calculated. This process was repeated for the baseline model, proposed model, and true data. These annual gains were also compared with long-term trading, where the stock was bought at the beginning of the year and sold at the end of the year.

Table 1: Shows the precision, recall, accuracy, and F1-scores for each class, and also shows the total accuracy of the baseline model's predicted classifications against test data.

Class	Precision	Recall	F1-Score	Accuracy
Sell (0)	0.310	0.562	0.400	0.5625
Buy (1)	0.238	0.714	0.357	0.7143
Hold (2)	0.943	0.764	0.844	0.7639
			Total Accuracy	0.748

Table 2: Shows the precision, recall, accuracy, and F1-scores for each class, and also shows the total accuracy of the proposed model's predicted classifications against test data.

Class	Precision	Recall	F1-Score	Accuracy
Sell (0)	0	0	0	0
Buy (1)	0	0	0	0
Hold (2)	0.878	1.000	0.935	1.00
			Total Accuracy	0.878

Table 3: Percentage gains from our models versus long-term trading and the maximum possible percent gain.

	Baseline Model	Proposed Model	Long-term Trading	Maximum Percent Gain
Total Percent Gain (%)	-16.36%	0%	50.65%	133.11%

The baseline model performed the worst in the financial evaluation, losing 16.36% of the original investment in one year. The proposed model neither lost nor gained in one year because it labeled all days as Hold position. The maximum gain possible in one year was 133.11%. In comparison, the Long-term trading analysis had a gain of 50.65% in one year (Table 3).

C. Limitations and Future Direction

Several limitations and future steps should be taken into consideration when moving forward. First, our data was limited to the scope of one stock symbol and our evaluation was during a year that pretty much only saw growth. This means that with a model that is trained on data with more fluctuations and the 2008 market crash, it would be much harder to evaluate correctly for the past year. Furthermore, this problem also has quite limited data when looking at one stock symbol. Our dataset only had around 5000 points and after down sampling to prevent class imbalance, we were left with only around 900 points of data. One future step could be to try oversampling. However, the challenge with this is the loss of the time series aspect. Several other models or learning types should be investigated. For example, a reinforcement learning strategy where the model is penalized for losing money could be implemented to see if performance improves. Our data could also be augmented to see if performance improves. First, window size in the algorithm used for labelling could be changed. And second, other data sources could be included since the stock market is heavily influenced by other factors such as the news. The implementation of an additional model that uses Natural Language Processing (NLP) could be added and would probably help with the prediction of faster fluctuations in market trends that make it "random."

REFERENCES

- [1] D. Nelson, A. M. Pereira, and Renato, "Stock market's price movement prediction with LSTM neural networks", 2017 International Joint Conference on Neural Networks (IJCNN), 2017. doi: 10.1109/IJCNN.2017.7966019 (accessed Feb. 27, 2021).
- [2] S. Hochreiter and J. Schmidhuber, "Long Short-Term Memory | Neural Computation," Neural Computation, 2020. <https://doi.org/10.1162/neco.1997.9.8.1735> (accessed Feb. 27, 2021).
- [3] NASDAQ. (2021, April 25). Advanced Micro Devices, Inc. (AMD). [Stock quote]. Retrieved from <https://finance.yahoo.com/quote/AMD>
- [4] O. B. Sezer and M. Ozbayoglu, "Algorithmic Financial Trading with Deep Convolutional Neural Networks: Time Series to Image Conversion Approach," Applied Soft Computing, vol. 70, pp. 525–538, Sep. 2018. <https://doi.org/10.1016/j.asoc.2018.04.024>